



# Convolutional Neural Network(CNN)




Our goal: recognize the image as boat




What is an image?

## Grey image


- A grid of numbers (*intensity values*)
- In computing a *matrix*



## Digit 8 in grey scale



# Color image



# Convolution

# Extracting features from images

- Extract features by applying a sliding window function across image, called as convolution
- Why sliding window?
  - Detect the same feature at different positions in the input image





How do you apply convolution across image?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image

1	0	1
0	1	0
1	0	1

Convolutional filter

How do you apply convolution across image?

1 <sub>x2</sub>	1 <sub>x0</sub>	1 <sub>x2</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x2</sub>	0 <sub>x0</sub>	1 <sub>x2</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Feature map

## Example1

Different filters (weights) reveal a different characteristics of the input.



$\ast 1/8$

0	1	0
1	4	1
0	1	0



## Example1

Different filters (weights) reveal a different characteristics of the input.



\*

0	-1	0
-1	4	-1
0	-1	0



Different filters (weights) reveal a different characteristics of the input.



\*

1	0	-1
2	0	-2
1	0	-1



## Example2: Feature maps

- The convolution of a filter1 (with red outline) slides over the input image to produce a feature map.
- The convolution of another filter2 (with the green outline), over the same image gives a different feature map


# Parameters for convolutional filters

# Parameters for convolutional filters

The size of the feature maps are controlled by following parameters of convolution filters:


- **Depth:** It corresponds to the number of filters we use for the convolution operation
- **Stride:** It is the number of pixels by which we slide our filter matrix over the input matrix
- **Padding:** Adding zeros around the border, so that we can apply the filter to bordering elements of our input image matrix

7 by 7 input, 3 by 3 filter with stride 1





7 by 7 input, 3 by 3 filter with stride 2



**3x3 output**

## Convolution for edge cases: zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

7 by 7 input, 3 by 3 filter with stride 1, pad with 1 pixel border


7 by 7 output

## Generalization: zero pad the border

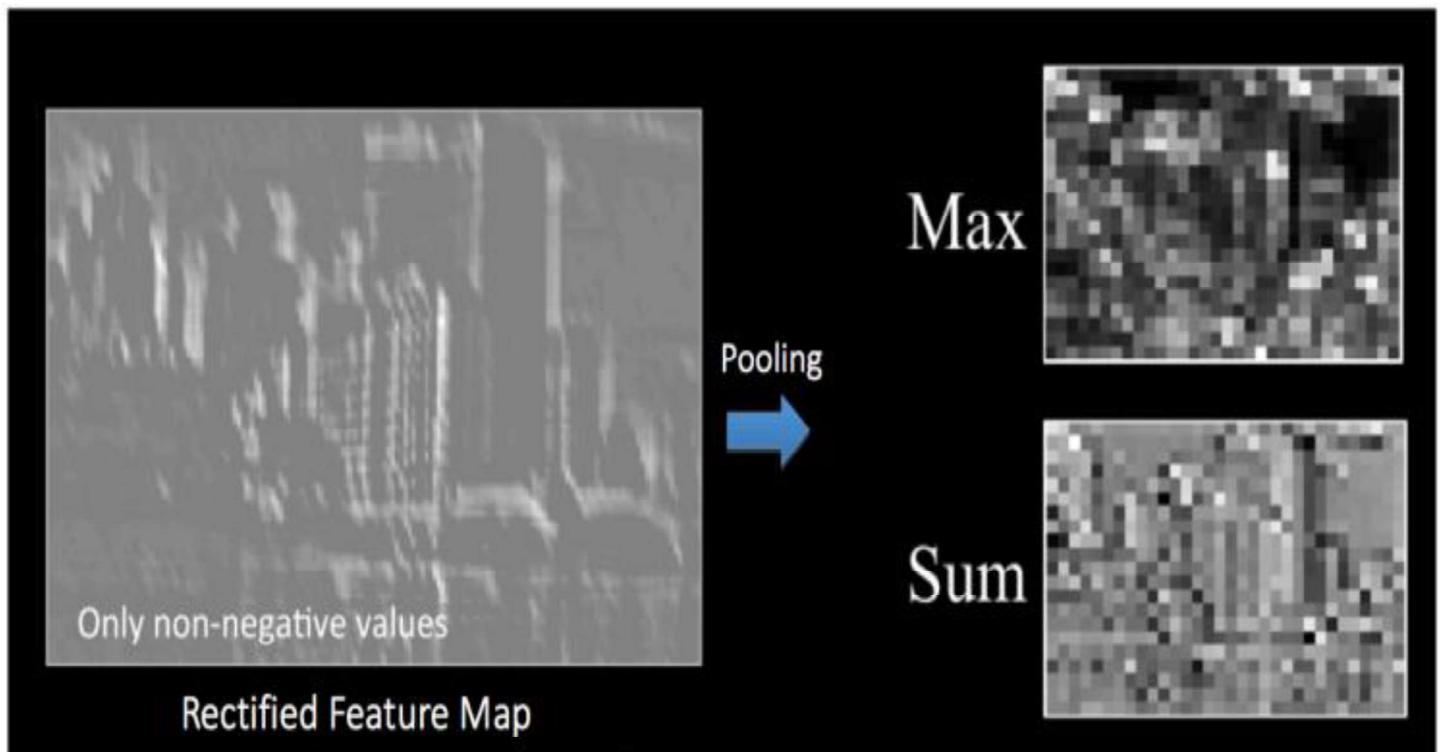
- In general, convolution layer will be having filters of size  $F$  by  $F$  with stride 1. The number of zero pad borders will be  $(F-1)/2$  to preserve size spatially
- E.g.,
  - $F = 3 \Rightarrow$  zero pad with 1
  - $F = 5 \Rightarrow$  zero pad with 2
  - $F = 7 \Rightarrow$  zero pad with 3

# Pooling/Subsampling

# Max pooling



# Max pooling




## What does pooling gives us?

- Makes the feature dimension smaller and fixed size
- Invariance to small transformations i.e., reduce the effect of noises and shift or distortion

Fully connected layers

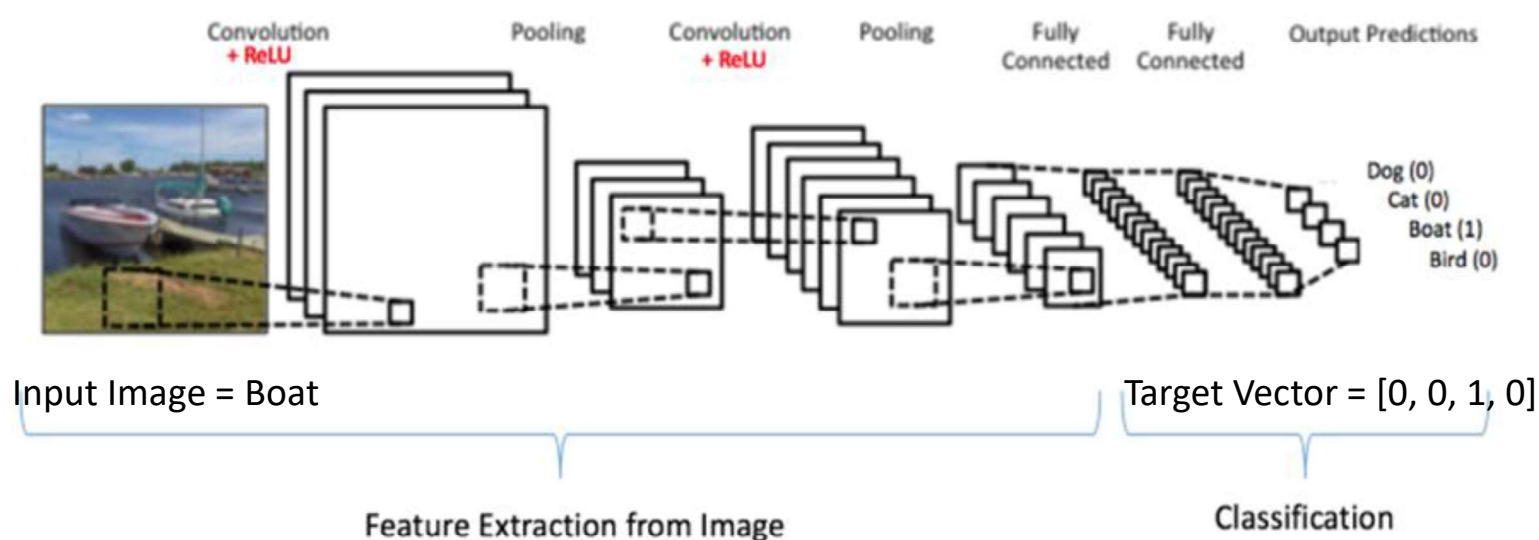


## FC layer



- The output from the convolutional and pooling layers represent high-level features of the input image
- The purpose of the Fully Connected layer is to use these features for classifying the input image into various classes based on the training dataset


# The final convolutional neural network




The Convolution + Pooling layers act as Feature Extractors from the input image while Fully Connected layer acts as a classifier.

# Perceptron view of CNN

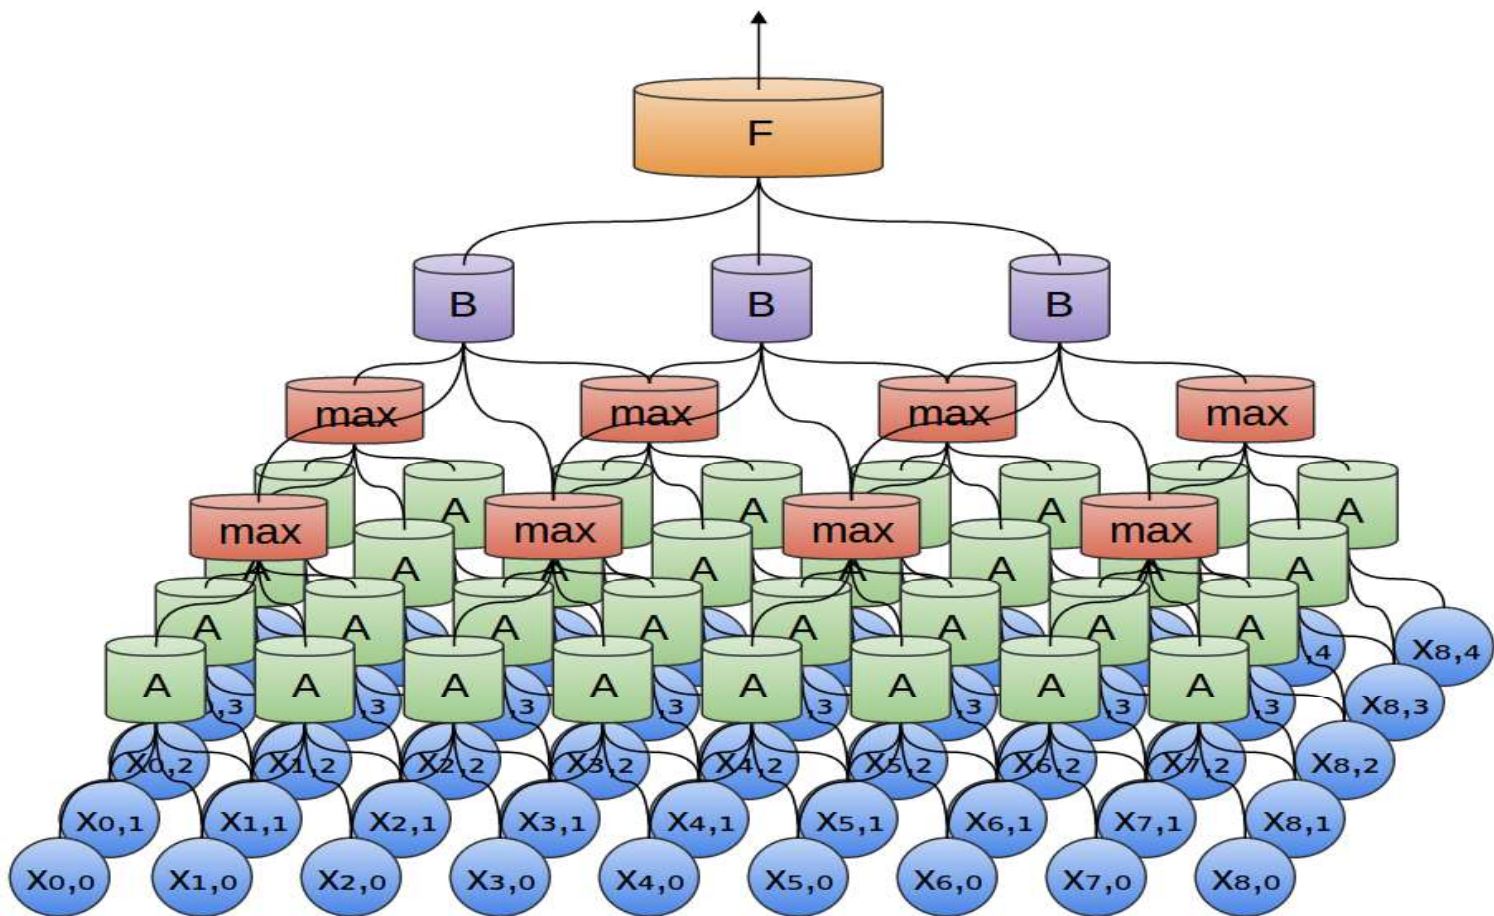
Convolution Filter = replicated  
perceptrons



Convolutional Filter = replicated  
perceptrons



Conv filter + Maxpooling + Conv filter +  
FC





# Channels & Activation maps

# 3D Data

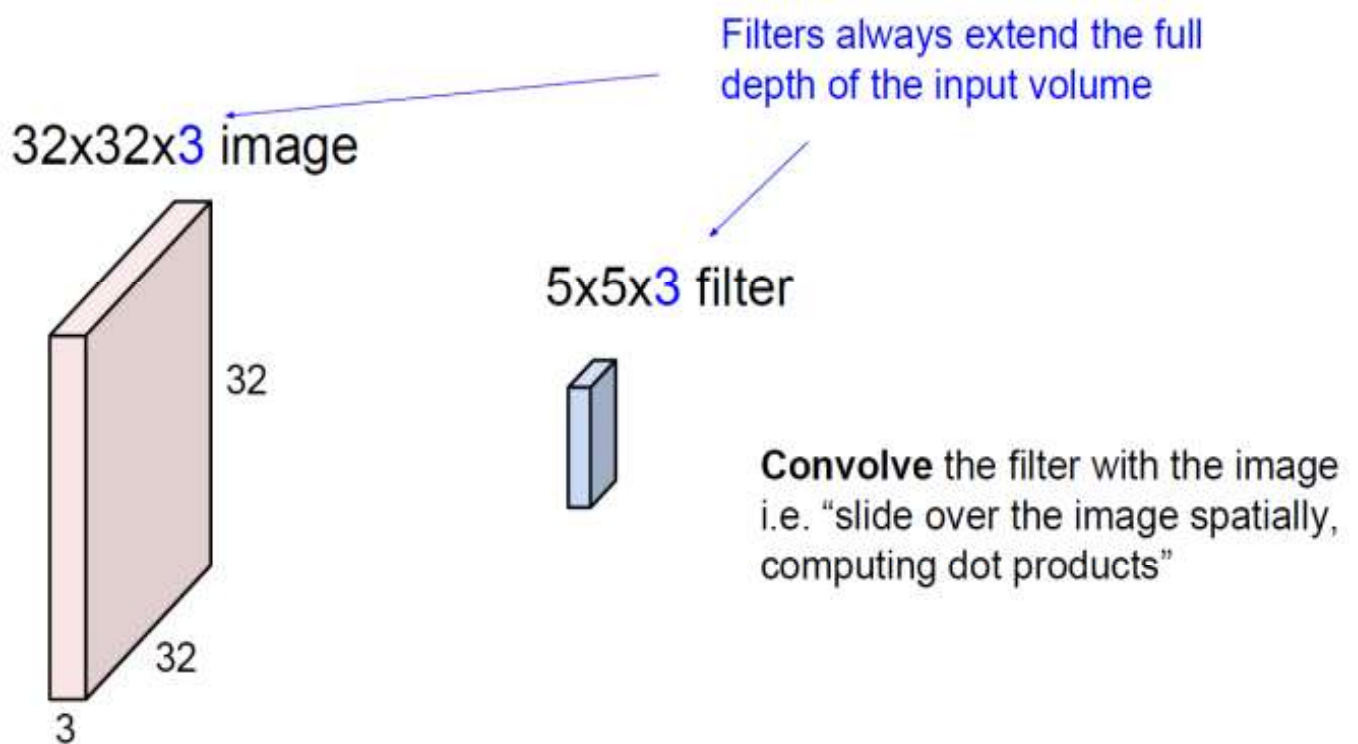
Data = 3D Tensor

- There is a vector of feature channels (e.g. RGB) at each spatial location (pixel).


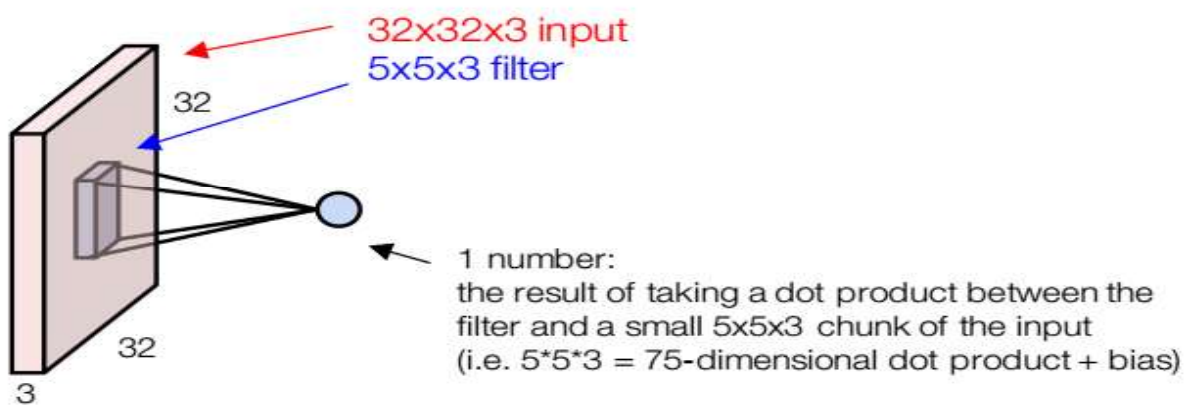





# Convolution with 3D Filters



# Convolution with 3D Filters




# Convolution with 3D Filters



# Convolution with 3D Filters

- Multiple filters produce multiple output channels
- For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get an output of size 28x28x6.

# Pooling across feature maps

- makes the representations smaller and more manageable
- operates over each activation map independently:

